

# DevOps Monitoring

Prometheus, Grafana & Observability for Modern Applications

# Monitoring Fundamentals

---

**Monitoring** is the practice of collecting, analyzing, and acting on data about your systems and applications to ensure reliability, performance, and availability.

## Why Monitoring in DevOps?

- **Proactive Issue Detection:** Identify problems before they impact users
- **Performance Optimization:** Understand system behavior and bottlenecks
- **Capacity Planning:** Make informed decisions about resource allocation
- **Compliance & SLA:** Meet service level agreements and regulatory requirements
- **Root Cause Analysis:** Quickly diagnose and resolve incidents
- **Business Intelligence:** Gain insights into user behavior and system usage

## Three Pillars of Observability

### Metrics

Numerical data representing system state over time (CPU, memory, request rate)

### Logs

Event details and contextual information about what happened in the system

### Traces

Distributed request tracking showing the path of requests through microservices

# Monitoring vs. Observability

Aspect	Monitoring	Observability
Focus	Known problems and predefined metrics	Understanding system behavior and unknown issues
Approach	Reactive - alerts when thresholds are breached	Proactive - enables exploration and investigation
Data	Aggregated metrics and dashboards	Rich context with metrics, logs, and traces
Questions	"Is the system working?"	"Why is the system behaving this way?"

## Importance of Observability

- **Complex Systems:** Modern microservices architectures are too complex for traditional monitoring
- **Unknown Unknowns:** Ability to investigate issues you didn't anticipate
- **Faster MTTR:** Reduce Mean Time To Resolution with better context
- **Developer Experience:** Enable developers to understand their applications in production
- **Business Impact:** Connect technical metrics to business outcomes

**Key Insight:** Observability is not just about collecting data - it's about making systems understandable and debuggable in production environments.

# Metrics, Logs, and Alerts

## Metrics: Numerical Data

**Metrics** are numerical measurements that represent the state of your system at a specific point in time, collected at regular intervals.

### Types of Metrics

#### Counter

Always increasing values  
(requests served, errors occurred)

#### Gauge

Values that can go up or down  
(CPU usage, memory consumption)

#### Histogram

Distribution of values (request duration, response sizes)

#### Summary

Quantiles and totals (95th percentile response time)

## Logs: Event Details

**Logs** are timestamped records of events that occurred in your system, providing context and details about what happened.


### Log Levels and Structure

```
# Structured Log Example
{
  "timestamp": "2023-10-23T10:30:00Z",
```

```
"level": "ERROR",
"service": "user-service",
"message": "Failed to authenticate user",
"user_id": "12345",
"error_code": "AUTH_FAILED",
"request_id": "req-abc123"
}
```

## Traces: Distributed Request Tracking

**Traces** show the journey of a request through multiple services, helping understand performance bottlenecks and dependencies.



## Alerts: Actionable Notifications

**Alerts** notify you when something requires attention, enabling quick response to issues before they impact users.

### Alert Types

Type	Description	Example
Threshold-based	Triggered when metric crosses predefined threshold	CPU usage > 80%
Anomaly-based	Triggered when behavior deviates from normal patterns	Request rate 3x higher than usual
Composite	Multiple conditions must be met	High CPU AND high memory

Type	Description	Example
Absence	Triggered when expected data is missing	No heartbeat for 5 minutes

## Importance of Alerts

- **Early Warning:** Detect issues before they become critical
- **Automation:** Enable automated responses to common issues
- **Escalation:** Ensure the right people are notified at the right time
- **Documentation:** Create audit trail of system events

**Alert Fatigue:** Too many alerts can desensitize teams. Focus on actionable alerts that require human intervention.

*Prepared By: Rashi Rana  
Corporate Trainer*

# SLA, SLO, SLI – Concepts

---

Service Level concepts provide a framework for measuring and communicating service reliability and performance expectations.

## Service Level Indicators (SLI)

### What are SLIs?

Quantitative measures of service behavior - the actual metrics you measure

- **Availability:** Percentage of successful requests
- **Latency:** Time to process requests (95th percentile)
- **Throughput:** Requests processed per second
- **Error Rate:** Percentage of failed requests

## Service Level Objectives (SLO)

### What are SLOs?

Target values for SLIs - internal goals for service performance

- **Availability SLO:** 99.9% uptime (8.76 hours downtime/year)
- **Latency SLO:** 95% of requests < 200ms
- **Error Rate SLO:** < 0.1% error rate

## Service Level Agreements (SLA)

## What are SLAs?

Contractual commitments to customers with consequences for not meeting them

- **Business Contract:** Legal agreement with customers
- **Penalties:** Credits or refunds for SLA breaches
- **Conservative:** Usually less strict than internal SLOs

## Relationship Between SLI, SLO, and SLA



## Example: E-commerce Website

Metric	SLI	SLO	SLA
Availability	Successful requests / Total requests	99.95% availability	99.9% availability
Latency	95th percentile response time	< 150ms for 95% of requests	< 200ms for 95% of requests
Error Rate	5xx errors / Total requests	< 0.05% error rate	< 0.1% error rate

## Error Budget

**Error Budget** =  $100\% - \text{SLO}$ . It represents how much unreliability you can tolerate while still meeting your objectives.

- **99.9% SLO = 0.1% Error Budget (43.8 minutes/month)**

- **Balance Innovation vs Reliability:** Spend error budget on new features
- **Freeze Deployments:** When error budget is exhausted

**Best Practice:** Start with SLIs that matter to users, set realistic SLOs based on current performance, and make SLAs more conservative than SLOs.

*Prepared By: Rashi Rana  
Corporate Trainer*

# Prometheus Architecture and Components

---

## What is Prometheus?

**Prometheus** is an open-source monitoring system that collects metrics from applications and infrastructure, stores them as time-series data, and provides powerful querying capabilities.

## Core Components

### Prometheus Server

Core component that scrapes, stores metrics and serves queries

### Exporters

Agents that expose metrics from systems (Node Exporter, MySQL Exporter)

### Pushgateway


For short-lived jobs that can't be scraped directly

### AlertManager

Handles alerts sent by Prometheus and routes them to receivers

## Prometheus Architecture on Kubernetes

Kubernetes Cluster



## How Prometheus Works (Pull Model)



### Pull Model Benefits

- **Centralized Control:** Prometheus decides what and when to scrape
- **Service Discovery:** Automatically finds new targets in Kubernetes
- **Health Monitoring:** Knows if targets are unreachable
- **Security:** Targets don't need to know about Prometheus location

## Grafana - Visualization Layer

**Grafana** connects to Prometheus as a data source and creates beautiful, interactive dashboards for monitoring data visualization.

### Key Grafana Features

- **Dashboard Creation:** Drag-and-drop interface for building dashboards
- **Multiple Data Sources:** Prometheus, InfluxDB, MySQL, etc.
- **Alerting:** Visual alerts with notification channels
- **Templating:** Dynamic dashboards with variables
- **Sharing:** Export/import dashboards as JSON

**Perfect Team:** Prometheus collects and stores metrics, Grafana visualizes them, and AlertManager handles notifications - creating a complete monitoring solution.

*Prepared By: Rashi Rana  
Corporate Trainer*

# Deploying Monitoring Stack with Helm

## Why Helm for Monitoring?

Helm simplifies Kubernetes deployments by packaging applications into charts. The kube-prometheus-stack chart includes Prometheus, Grafana, and AlertManager pre-configured.

## kube-prometheus-stack Chart

### Complete Stack

Prometheus, Grafana,  
AlertManager in one chart

### Pre-configured

Ready-to-use dashboards and  
alerts

### Kubernetes Native

ServiceMonitors,  
PrometheusRules CRDs included

### Customizable

Override values for specific  
requirements

## Step-by-Step Deployment

1

### Add Helm Repository

```
helm repo add prometheus-community https://prometheus-  
community.github.io/helm-charts  
helm repo update
```

2

### Create Namespace

```
kubectl create namespace monitoring
```

3

### Install the Stack

```
helm install prometheus prometheus-community/kube-prometheus-stack \
--namespace monitoring \
--set grafana.adminPassword=admin123
```

4

### Verify Installation

```
kubectl get pods -n monitoring
kubectl get svc -n monitoring
```

## What Gets Deployed

Component	Service Type	Default Port	Purpose
<b>Prometheus</b>	ClusterIP	9090	Metrics collection and storage
<b>Grafana</b>	ClusterIP	80	Dashboard and visualization
<b>AlertManager</b>	ClusterIP	9093	Alert routing and notifications
<b>Node Exporter</b>	DaemonSet	9100	Node-level metrics

## Accessing the Services

```
# Access Prometheus UI
kubectl port-forward -n monitoring svc/prometheus-kube-prometheus-prometheus 9090:9090
# Open: http://localhost:9090

# Access Grafana UI
kubectl port-forward -n monitoring svc/prometheus-grafana 3000:80
# Open: http://localhost:3000 (admin/admin123)

# Access AlertManager UI
```

```
kubectl port-forward -n monitoring svc/prometheus-kube-prometheus-
alertmanager 9093:9093
# Open: http://localhost:9093
```

## Custom Values Configuration

```
# values.yaml
prometheus:
  prometheusSpec:
    retention: 30d
    storageSpec:
      volumeClaimTemplate:
        spec:
          storageClassName: standard
          accessModes: ["ReadWriteOnce"]
        resources:
          requests:
            storage: 50Gi

  grafana:
    adminPassword: mySecurePassword
    service:
      type: LoadBalancer
    persistence:
      enabled: true
      size: 10Gi

  alertmanager:
    alertmanagerSpec:
      storage:
        volumeClaimTemplate:
          spec:
            storageClassName: standard
            accessModes: ["ReadWriteOnce"]
          resources:
            requests:
              storage: 10Gi
```

## Deploy with Custom Values

```
# Install with custom configuration
helm install prometheus prometheus-community/kube-prometheus-stack
\ 
--namespace monitoring \
--values values.yaml
```

```
# Upgrade existing installation
helm upgrade prometheus prometheus-community/kube-prometheus-stack
\
--namespace monitoring \
--values values.yaml
```

**Production Ready:** The Helm chart provides a production-ready monitoring stack with persistent storage, proper RBAC, and pre-configured dashboards!


*Prepared By: Rashi Rana  
Corporate Trainer*

# AlertManager Architecture and Configuration

## What is AlertManager?

**AlertManager** receives alerts from Prometheus, groups them, and routes them to the correct notification channels like email, Slack, or PagerDuty.

## AlertManager Architecture



## Key AlertManager Features

**Grouping**

**Routing**

Combines similar alerts to reduce notification spam

Sends alerts to different teams based on labels


### Inhibition

Suppresses alerts when higher priority ones are active

### Silencing

Temporarily mute alerts during planned maintenance

## Alert Flow in Kubernetes



## Basic AlertManager Configuration

```
# alertmanager.yml
global:
  smtp_smarthost: 'localhost:587'
  smtp_from: 'alerts@company.com'

route:
  group_by: ['alertname', 'cluster']
  group_wait: 10s
  group_interval: 10s
  repeat_interval: 1h
  receiver: 'web.hook'
  routes:
    - match:
        severity: critical
        receiver: 'critical-alerts'

receivers:
  - name: 'web.hook'
webhook_configs:
  - url: 'http://127.0.0.1:5001/'
  - name: 'critical-alerts'
email_configs:
  - to: 'admin@company.com'
```

```

subject: 'Critical Alert: {{ .GroupLabels.alertname }}'
body: 'Alert details: {{ range .Alerts }}{{ .Annotations.summary }}{{ end }}'

```

## Alert States

State	Description	Action
<b>Inactive</b>	Alert condition is not met	No action taken
<b>Pending</b>	Condition met but waiting for 'for' duration	Monitoring, not yet firing
<b>Firing</b>	Condition met for specified duration	Alert sent to AlertManager

## Creating Alert Rules

```

# prometheus-rules.yaml
apiVersion: monitoring.coreos.com/v1
kind: PrometheusRule
metadata:
  name: basic-alerts
  namespace: monitoring
spec:
  groups:
    - name: basic.rules
      rules:
        - alert: HighCPUUsage
          expr: 100 - (avg(irate(node_cpu_seconds_total{mode="idle"} [5m])) * 100) > 80
          for: 2m
          labels:
            severity: warning
          annotations:
            summary: "High CPU usage on {{ $labels.instance }}"
            description: "CPU usage is above 80% for more than 2 minutes"
        - alert: PodCrashLooping
          expr: rate(kube_pod_container_status_restarts_total[15m]) > 0
          for: 5m
          labels:
            severity: critical
          annotations:
            summary: "Pod {{ $labels.pod }} is crash looping"

```

```
description: "Pod has restarted {{ $value }} times in the last 15  
minutes"
```

**Smart Alerting:** AlertManager ensures you get the right alerts to the right people at the right time, without overwhelming your team with noise.

*Prepared By: Rashi Rana  
Corporate Trainer*

# Alertmanager and SLA Thresholds

## Introduction to Alertmanager

**Alertmanager** handles alerts sent by Prometheus server, taking care of deduplicating, grouping, and routing them to correct receivers like email, Slack, or PagerDuty.

## Alertmanager Features

### Grouping

Group similar alerts to reduce noise

### Inhibition

Suppress alerts when other alerts are firing

### Silencing

Temporarily mute alerts during maintenance

### Routing

Send alerts to different teams based on labels

## Creating Alerts in Prometheus

```
# prometheus-rules.yaml
groups:
- name: example-alerts
rules:
- alert: HighCPUUsage
expr: 100 - (avg by(instance)
(irate(node_cpu_seconds_total{mode="idle"} [5m])) * 100) > 80
for: 2m
labels:
severity: warning
```

```

team: infrastructure
annotations:
summary: "High CPU usage detected"
description: "CPU usage is above 80% for more than 2 minutes on {{ $labels.instance }}"
  - alert: HighMemoryUsage
expr: (node_memory_MemTotal_bytes - node_memory_MemAvailable_bytes) / node_memory_MemTotal_bytes * 100 > 85
for: 5m
labels:
severity: critical
team: infrastructure
annotations:
summary: "High memory usage detected"
description: "Memory usage is above 85% for more than 5 minutes on {{ $labels.instance }}"

```

## Example Alert Rules

Alert	Expression	Threshold
CPU Alert	100 - (avg(irate(node_cpu_seconds_total{mode="idle"}[5m])) * 100)	> 80%
Memory Alert	(1 - node_memory_MemAvailable_bytes/node_memory_MemTotal_bytes) * 100	> 85%
Disk Alert	(1 - node_filesystem_avail_bytes/node_filesystem_size_bytes) * 100	> 90%
Pod Restart	increase(kube_pod_container_status_restarts_total[1h])	> 5

## SLA Thresholds and Use

Set alert thresholds based on your SLOs to ensure you're notified before SLA breaches occur.

### SLA-based Alert Strategy

- **Error Budget Alerts:** Alert when 50% of error budget is consumed
- **Burn Rate Alerts:** Alert on fast error budget consumption
- **Availability Alerts:** Alert when availability drops below SLO
- **Latency Alerts:** Alert when response times exceed SLO

```
# SLA-based alert example
- alert: ErrorBudgetBurn
expr: (
  1 - (
    sum(rate(http_requests_total{code!~"5.."}[5m])) /
    sum(rate(http_requests_total[5m]))
  )
) > 0.001 # 0.1% error rate (SLO breach)
for: 2m
labels:
severity: critical
annotations:
summary: "SLA breach: Error rate above 0.1%"
```

## Creating Alerts from Grafana GUI

Grafana provides a user-friendly interface to create alerts directly from dashboards without writing YAML files.

### Steps to Create GUI Alerts

- 1 **Open Dashboard Panel:** Click on panel title → Edit
- 2 **Go to Alert Tab:** Click "Alert" tab in panel editor
- 3 **Create Alert Rule:** Click "Create Alert" button
- 4 **Set Conditions:** Define query, condition, and threshold
- 5 **Configure Notifications:** Select notification channels
- 6 **Save Alert:** Test and save the alert rule

### GUI Alert Configuration

Query

Condition

Use existing panel query or  
create new one

IS ABOVE, IS BELOW, IS  
OUTSIDE RANGE

## Evaluation

Set frequency and duration for  
alert evaluation

## Notifications

Choose Slack, email, webhook  
channels

### Example GUI Alert Setup

```
# GUI Alert Configuration Example:  
Query: avg(cpu_usage_percent)  
Condition: IS ABOVE 80  
Evaluation: every 1m for 2m  
Notification: #alerts-channel  
Message: "CPU usage is {{ $value }}% on {{ $labels.instance }}"
```

**Alert Tuning:** Start with conservative thresholds and adjust based on false positive rates and actual incidents.

*Prepared By: Rashi Rana  
Corporate Trainer*

# Hands-on Lab: Complete Monitoring Setup

---

## Lab Objective

Deploy a complete monitoring stack with Prometheus, Grafana, and AlertManager on Kubernetes using Helm, then configure alerts and dashboards.

## Lab Steps Overview

### Step 1

Deploy monitoring stack with Helm

### Step 2

Configure Slack integration

### Step 3

Create custom alerts

### Step 4

Build Grafana dashboard

## Step 1: Deploy Monitoring Stack

```
# Add Helm repository
helm repo add prometheus-community https://prometheus-
community.github.io/helm-charts
helm repo update

# Create namespace
```

```
kubectl create namespace monitoring

# Install complete stack
helm install prometheus prometheus-community/kube-prometheus-stack \
\ --namespace monitoring \
--set grafana.adminPassword=admin123 \
--set prometheus.prometheusSpec.retention=30d

# Verify installation
kubectl get pods -n monitoring
kubectl get svc -n monitoring
```

## Step 2: Access Services

```
# Access Prometheus (Terminal 1)
kubectl port-forward -n monitoring svc/prometheus-kube-prometheus-
prometheus 9090:9090
# Open: http://localhost:9090

# Access Grafana (Terminal 2)
kubectl port-forward -n monitoring svc/prometheus-grafana 3000:80
# Open: http://localhost:3000 (admin/admin123)

# Access AlertManager (Terminal 3)
kubectl port-forward -n monitoring svc/prometheus-kube-prometheus-
alertmanager 9093:9093
# Open: http://localhost:9093
```

## Step 3: Configure Slack Alerts

```
# Create alertmanager-config.yaml
apiVersion: v1
kind: Secret
metadata:
name: alertmanager-main
namespace: monitoring
stringData:
alertmanager.yml: |
global:
slack_api_url: 'https://hooks.slack.com/services/YOUR/SLACK/WEBHOOK'
route:
group_by: ['alertname']
group_wait: 10s
```

```
group_interval: 10s
repeat_interval: 1h
receiver: 'slack-notifications'
receivers:
- name: 'slack-notifications'
slack_configs:
- channel: '#monitoring'
title: 'Alert: {{ .GroupLabels.alertname }}'
text: '{{ range .Alerts }}{{ .Annotations.summary }}{{ end }}'
```

```
# Apply configuration
kubectl apply -f alertmanager-config.yaml

# Restart AlertManager
kubectl rollout restart statefulset/alertmanager-prometheus-kube-
prometheus-alertmanager -n monitoring
```

## Step 4: Create Custom Alert

```
# Create custom-alert.yaml
apiVersion: monitoring.coreos.com/v1
kind: PrometheusRule
metadata:
name: custom-alerts
namespace: monitoring
labels:
prometheus: kube-prometheus
role: alert-rules
spec:
groups:
- name: custom.rules
rules:
- alert: HighCPUUsage
expr: 100 - (avg(irate(node_cpu_seconds_total{mode="idle"} [5m])) *
100) > 80
for: 2m
labels:
severity: warning
annotations:
summary: "High CPU usage detected"
description: "CPU usage is {{ $value }}% for more than 2 minutes"
```

```
# Apply custom alert
kubectl apply -f custom-alert.yaml
```

```
# Check if alert is loaded  
kubectl get prometheusrules -n monitoring
```

## Step 5: Create Grafana Dashboard

- 1** **Login to Grafana:** <http://localhost:3000> (admin/admin123)
- 2** **Create Dashboard:** Click "+" → "Dashboard"
- 3** **Add Panel:** Click "Add new panel"
- 4** **Configure Query:**

```
# CPU Usage Query  
100 - (avg(irate(node_cpu_seconds_total{mode="idle"}[5m]))  
* 100)  
  
# Memory Usage Query  
(1 - (node_memory_MemAvailable_bytes /  
node_memory_MemTotal_bytes)) * 100
```

- 5** **Save Dashboard:** Give it a name and save

**Lab Complete!** You now have a fully functional monitoring stack with Prometheus collecting metrics, Grafana visualizing them, and AlertManager sending notifications to Slack.

*Prepared By: Rashi Rana  
Corporate Trainer*

# Summary and Key Takeaways

---

## What We Learned

You now understand how to implement a complete monitoring solution using Prometheus, Grafana, and AlertManager on Kubernetes with Helm.

## Core Concepts Mastered

### Monitoring Fundamentals

Three pillars: Metrics, Logs, Traces

### SLA/SLO/SLI

Service level concepts for reliability

### Prometheus Architecture

Pull-based metrics collection system


### Grafana Dashboards

Visualization and alerting platform

## Technical Skills Gained

- **Helm Deployment:** Deploy monitoring stack with kube-prometheus-stack chart
- **Alert Configuration:** Create PrometheusRules and configure AlertManager
- **Slack Integration:** Set up notification channels for team collaboration
- **Dashboard Creation:** Build custom Grafana dashboards with PromQL queries
- **Kubernetes Monitoring:** Monitor pods, services, and cluster resources

## Architecture Understanding



## Production Best Practices

- **Persistent Storage:** Configure storage for metrics retention
- **Resource Limits:** Set appropriate CPU/memory limits
- **High Availability:** Deploy multiple replicas for critical components
- **Security:** Enable authentication and RBAC
- **Backup:** Regular backup of Grafana dashboards and configurations

## Next Steps for Learning

- 1 **Advanced PromQL:** Learn complex queries and functions
- 2 **Custom Exporters:** Create application-specific metrics exporters
- 3 **Distributed Tracing:** Implement Jaeger or Zipkin for microservices
- 4 **Log Management:** Add ELK/EFK stack for centralized logging
- 5 **Multi-cluster Monitoring:** Federated Prometheus setup

## Common Commands Reference

```
# Deploy monitoring stack
helm install prometheus prometheus-community/kube-prometheus-stack
-n monitoring

# Access services
kubectl port-forward -n monitoring svc/prometheus-kube-prometheus-
prometheus 9090:9090
kubectl port-forward -n monitoring svc/prometheus-grafana 3000:80
kubectl port-forward -n monitoring svc/prometheus-kube-prometheus-
```

```
alertmanager 9093:9093

# Create custom alerts
kubectl apply -f prometheus-rules.yaml

# Check alert rules
kubectl get prometheusrules -n monitoring
```

## Key Resources

- **Prometheus Documentation:** <https://prometheus.io/docs/>
- **Grafana Documentation:** <https://grafana.com/docs/>
- **Helm Charts:** <https://prometheus-community.github.io/helm-charts/>
- **PromQL Guide:**  
<https://prometheus.io/docs/prometheus/latest/querying/basics/>

**Congratulations!** You're now equipped to implement comprehensive monitoring solutions for modern applications and infrastructure. Start monitoring your applications today!

*Prepared By: Rashi Rana  
Corporate Trainer*